

INTERIM ARCHITECTURE REPORT (TEXT PLACEHOLDER)

Executive summary

This interim report documents the current design and implementation progress of the **“Automaton Auditor” Digital Courtroom**. The codebase includes typed state models ([Evidence](#), [JudicialOpinion](#), [AgentState](#)), sandboxed git tooling, and a LangGraph [StateGraph](#) wiring the Detective layer in parallel with a synchronization node.

Architecture decisions so far (with trade-offs)

Pydantic + [TypedDict](#) over plain dicts

- **Decision:** Use Pydantic [BaseModel](#) for [Evidence](#) and [JudicialOpinion](#) to ensure strongly-typed, validated objects exchanged between nodes. Keep [AgentState](#) as a [TypedDict](#) with [Annotated](#) reducers (e.g., `operator.add`, `operator.ior`) so parallel branches can safely merge lists and dicts.
 - **Trade-offs:** Pydantic adds runtime validation overhead and stronger coupling to the schema. In this governance/audit setting the clarity and safety provided by types outweigh the small performance cost.
 - **Alternatives considered:**
 - Pure dict-based state — rejected (too brittle for parallel reducers, hard to validate at graph boundaries).
 - Single monolithic Pydantic model — rejected (would complicate LangGraph integration and create tightly coupled schemas).
-

AST-based analysis over regex scanning

- **Decision:** Use Python's `ast` module (and possibly `tree-sitter`) for repo forensics rather than regex-based searches to understand graph wiring and state structures.
 - **Trade-offs:** AST parsing is more complex and heavier than simple string matching, but robust to formatting changes and avoids false positives (e.g., commented mentions of `StateGraph`).
 - **Rejected alternative:** Regex scanning — considered for speed/simplicity but rejected because the rubric requires structural verification rather than mere token presence.
-

Sandboxed git tooling

- **Decision:** `src/tools/repo_tools.py` clones repositories into a `tempfile.TemporaryDirectory` and runs `git log --oneline --reverse` to build a commit narrative; it never touches the live working tree.
 - **Trade-offs:** Temporary directories require careful lifetime management and slightly more code, but they reduce the blast radius of cloning untrusted repos.
 - **Rejected alternative:** Cloning directly into the workspace (e.g., `os.system("git clone ...")`) — rejected as a security liability.
-

RAG-lite PDF forensics

- **Decision:** `src/tools/doc_tools.py` ingests PDFs once, chunks/indexes content, and supports targeted queries for concepts such as Dialectical Synthesis, Fan-In/Fan-Out, and Metacognition — rather than dumping the entire report into a single LLM prompt.
 - **Trade-offs:** RAG-lite adds complexity (chunking, indexing, targeted queries) but keeps context windows small and allows citing specific sentences as Evidence.
 - **Rejected alternative:** Naive single-prompt summarizer — rejected because it's hard to ground forensic claims to exact text spans.
-

StateGraph flow (Detective focus)

1. **Entry:** `context_builder` loads rubric dimensions into state.
 2. **Fan-out Detectives (parallel):**
 - o `repo_investigator` — Code Detective (git, AST, graph, state)
 - o `doc_analyst` — Paperwork Detective (PDF and textual forensics)
 - o `vision_inspector` — Diagram Detective (image/diagram analysis; currently a structural placeholder)
 3. **Fan-in:** `evidence_aggregator` synchronizes inputs. Because `evidences` uses an `operator.ior` reducer, each Detective can contribute evidence without overwriting others.
 4. **Judicial fan-out:** Evidence is reviewed by multiple judges (parallel), producing `JudicialOpinion` lists.
 5. **Chief Justice:** Deterministic synthesis of opinions into an `AuditReport` (Markdown).
-

Known gaps, risks, and forward plan

RepoInvestigator

- **Planned:** Implement AST parsing to validate `AgentState` structure, reducers, and StateGraph fan-out/fan-in wiring. Capture the exact graph definition block as Evidence.
- **Risks:** AST visitors might miss dynamically constructed graphs or unusual `add_edge` patterns, causing false negatives.
- **Mitigation:** Start with a narrow set of supported patterns and document assumptions in the Evidence rationale.

DocAnalyst

- **Planned:** Implement robust PDF parsing and file-path extraction; cross-reference claims with repository contents to flag hallucinated paths.
- **Risks:** OCR/parse errors or unusual formatting can hide paths and phrases, causing under-reporting.
- **Mitigation:** Keep raw extracted text snippets in Evidence and enable manual inspection during debugging.

VisionInspector

- **Planned:** Add image extraction and multimodal analysis to classify diagrams and verify that parallel Detectives/Judges and Chief Justice synthesis are shown rather than a linear pipeline.
- **Risks:** Small, dense, or low-quality diagrams may be misclassified.
- **Mitigation:** Treat vision findings as complementary and explicitly record absence of diagrams.

Judicial layer & Chief Justice

- **Planned:** Replace placeholders with LLM-backed judge personas that output [JudicialOpinion](#) (structured). Implement deterministic ChiefJustice rules (security override, fact supremacy, dissent requirement) to synthesize scores into a Markdown [AuditReport](#).
 - **Risks:** Persona collusion (near-identical opinions) or hallucinated citations.
 - **Mitigation:** Enforce `.with_structured_output(JudicialOpinion)` with strict validation and retries; design prompts to emphasize philosophical differences.
-

Sequencing & prioritization

1. **Stabilize forensic foundation** — AST-based RepoInvestigator + robust PDF ingestion (foundation for all higher judgments).

2. **Harden judicial personas** — structured LLM judges, tune prompts to produce meaningfully divergent, grounded opinions.
 3. **Implement Chief Justice rules** — encode overrides, fact supremacy and dissent checks as deterministic Python logic; wire score-variance checks.
 4. **Enhance diagrams & vision** — add StateGraph diagrams and VisionInspector after textual evidence and judicial reasoning are reliable.
-

StateGraph architecture diagram (Mermaid)

